

Homework 04

Jiarui Huang

202283890036

IoT

Exercise 5.1 In this exercise we look at memory locality properties of matrix computation. The following code is written in C, where element within the same row are stored contiguously. Assume each word is a 32-bit integer.

```
For(I=0; I<8; I++)  
    For(J=0; J<8000; J++)  
        A[I][J]=B[I][J]+A[I][J];
```

5.1.1 How many 32-bit integers can be stored in a 16-byte cache block?

5.1.2 Which variable references exhibit temporal locality?

5.1.3 Which variable references exhibit spatial locality?

SOLUTION:

5.1.1 4

5.1.2 I, J and A[I][J]

5.1.3 A[I][J] B[I][J]

Exercise 5.2 Caches are important to providing a high-performance memory hierarchy to processors. Below is a list of 32-bit memory address references, given a word addresses.

0x03, 0xb4, 0x2b, 0x02, 0xbf, 0x58, 0xbe, 0x0e, 0xb5, 0x2c, 0xba, 0xfd

5.2.1 For each of these references, identify the binary address, the tag, and the index given a direct-mapped cache with 16 one-word blocks. Also list if each reference is a hit or a miss, assuming the cache is initially empty

5.2.2 For each of these references, identify the binary address, the tag, and the index given a direct-mapped cache with two-word blocks and a total size of 8 blocks. Also list if each reference is a hit or a miss, assuming the cache is initially empty.

5.2.3 You are asked to optimize a cache design for the given references. There are three direct-mapped cache designs possible, all with a total of 8 words of data:

- C1 has 1-word blocks,
- C2 has 2-word blocks, and

- C3 has 4-word blocks

SOLUTION:

5.2.1

| Address | Binary address | Tag | Index | Hit or miss |
|---------|----------------|------|-------|-------------|
| 0x03 | 0000 0011 | 0000 | 0000 | miss |
| 0xb4 | 1011 0100 | 1011 | 0100 | miss |
| 0x2b | 0010 1011 | 0010 | 1011 | miss |
| 0x02 | 0000 0010 | 0000 | 0010 | miss |
| 0xbf | 1011 1111 | 1011 | 1111 | miss |
| 0x58 | 0101 1000 | 0101 | 1000 | miss |
| 0xbe | 1011 1110 | 1011 | 1110 | miss |
| 0x0e | 0000 1110 | 0000 | 1110 | miss |
| 0xb5 | 1011 0101 | 1011 | 0101 | miss |
| 0x2c | 0010 1100 | 0010 | 1100 | miss |
| 0xba | 1011 1010 | 1011 | 1010 | miss |
| 0xfd | 1111 1101 | 1111 | 1101 | miss |

5.2.2

| Address | Binary address | Tag | Index | Store block | Hit or miss |
|---------|----------------|-------|-------|-------------|-------------|
| 0x03 | 0000 0011 | 00000 | 01 | 1 | miss |
| 0xb4 | 1011 0100 | 10110 | 10 | 0 | miss |
| 0x2b | 0010 1011 | 00101 | 01 | 1 | miss |
| 0x02 | 0000 0010 | 00000 | 01 | 0 | hit |
| 0xbf | 1011 1111 | 10111 | 11 | 1 | miss |
| 0x58 | 0101 1000 | 10111 | 00 | 0 | miss |
| 0xbe | 1011 1110 | 10111 | 11 | 0 | hit |
| 0x0e | 0000 1110 | 00001 | 11 | 0 | miss |
| 0xb5 | 1011 0101 | 10110 | 10 | 1 | hit |
| 0x2c | 0010 1100 | 00101 | 10 | 0 | miss |
| 0xba | 1011 1010 | 10111 | 01 | 0 | miss |
| 0xfd | 1111 1101 | 11111 | 10 | 1 | miss |

5.2.3

| Address | Binary address | Cache 1 (1-word) | | | Cache 2 (2-word) | | | Cache 3 (4-word) | | |
|---------|----------------|------------------|-------|----------|------------------|-------|----------|------------------|-------|----------|
| | | Tag | Index | Hit/miss | Tag | Index | Hit/miss | Tag | Index | Hit/miss |
| 0x03 | 0000 0011 | 00000 | 000 | miss | 00000 | 01 | miss | 00000 | 0 | miss |
| 0xb4 | 1011 0100 | 10110 | 100 | miss | 10110 | 10 | miss | 10110 | 0 | miss |
| 0x2b | 0010 1011 | 00101 | 011 | miss | 00101 | 01 | miss | 00101 | 1 | miss |
| 0x02 | 0000 0010 | 00000 | 010 | miss | 00000 | 01 | hit | 00000 | 0 | hit |

| | | | | | | | | | | |
|------|-----------|-------|-----|------|-------|----|------|-------|---|------|
| 0xbf | 1011 1111 | 10111 | 111 | miss | 10111 | 11 | miss | 10111 | 1 | miss |
| 0x58 | 0101 1000 | 10111 | 000 | miss | 10111 | 00 | miss | 10111 | 1 | miss |
| 0xbe | 1011 1110 | 10111 | 110 | miss | 10111 | 11 | hit | 10111 | 1 | hit |
| 0x0e | 0000 1110 | 00001 | 110 | miss | 00001 | 11 | miss | 00001 | 1 | miss |
| 0xb5 | 1011 0101 | 10110 | 101 | miss | 10110 | 10 | hit | 10110 | 0 | hit |
| 0x2c | 0010 1100 | 00101 | 100 | miss | 00101 | 10 | miss | 00101 | 1 | miss |
| 0xba | 1011 1010 | 10111 | 010 | miss | 10111 | 01 | miss | 10111 | 1 | miss |
| 0xfd | 1111 1101 | 11111 | 101 | miss | 11111 | 10 | miss | 11111 | 1 | miss |

Exercise 5.3 By convention, a cache is named according to the amount of data it contains (i.e., a 4 KiB cache can hold 4 KiB of data); however, caches also require SRAM to store metadata such as tags and valid bits. For this exercise, you will examine how a cache's configuration affects the total amount of SRAM needed to implement it, as well as the performance of the cache. For all parts, assume that the caches are byte addressable, and that addresses and words are 64 bits.

- 5.3.1** Calculate the total number of bits required to implement a 32 KiB cache with two-word blocks.
- 5.3.2** Calculate the total number of bits required to implement a 64 KiB cache with 16-word blocks. How much bigger is this cache than the 32 KiB cache described in Exercise 5.3.1? (Notice that, by changing the block size, we doubled the amount of data without doubling the total size of the cache.)
- 5.3.2** Explain why this 64 KiB cache, despite its larger data size, might provide slower performance than the first cache.

SOLUTION:

- 5.3.1** The total cache contains is :

$$32\text{KiB} = 32 \times 2^{10} = 2^{15} \text{ bytes}$$

Of each block has two words, thus the number of lines is:

$$\log_2(2^{15} \div (2^3 \times 2)) = \log_2 2^{11} = 11$$

The 64-bits addresses are divided into : 1-bit for block offset, 11-bits for index offset, 3-bit for word offset and 49-bits for tag ($64 - 1 - 3 - 11 = 49$),

The cache is composed with:

$$2^{15} \times 8 + 2^{11} \times (49 + 1) = 364544 \text{ bits}$$

- 5.3.2** The total cache contains is:

$$64\text{KiB} = 64 \times 2^{10} = 2^{16} \text{ bytes}$$

With each block is 16 words blocks, for each block, the number of data can store is:

$$16 \times 2^3 = 2^7 \text{ bytes}$$

The lines of cache needs is:

$$\log_2(2^{16} \div 2^7) = \log_2 2^9 = 9 \text{ lines}$$

The 64-bits addresses are divided into: 4-bits for block offset, 9-bits for index offset

t, 3-bit for word offset and 48-bits for tag ($64-4-3-9=48$),

The cache is composed with:

$$2^{16} \times 8 + 2^9 \times (1 + 48) = 549376 \text{ bits}$$

- 5.3.3** The larger block size may require an increased hit time and an increased miss penalty than the original cache. The fewer number of blocks may cause a higher conflict miss rate than the original cache.

Exercise 5.8 Media applications that play audio or video files are part of a class of workloads called “streaming” workloads (i.e., they bring in large amounts of data but do not reuse much of it). Consider a video streaming workload that accesses a 512 KiB working set sequentially with the following address stream:
0, 1, 2, 3, 4, 5, 6, 7, 8, 9 ...

- 5.8.1** Assume a 64 KiB direct-mapped cache with a 32-byte block. What is the miss rate for the address stream above? How is this miss rate sensitive to the size of the cache or the working set? How would you categorize the misses this workload is experiencing, based on the 3C model?
- 5.8.2** Re-compute the miss rate when the cache block size is 16 bytes, 64 bytes, and 128 bytes. What kind of locality is this workload exploiting?
- 5.8.3** “Prefetching” is a technique that leverages predictable address patterns to speculatively bring in additional cache blocks when a particular cache block is accessed. One example of prefetching is a stream buffer that prefetches sequentially adjacent cache blocks into a separate buffer when a particular cache block is brought in. If the data is found in the prefetch buffer, it is considered as a hit and moved into the cache and the next cache block is prefetched. Assume a two-entry stream buffer and assume that the cache latency is such that a cache block can be loaded before the computation on the previous cache block is completed. What is the miss rate for the address stream above?

SOLUTION:

- 5.8.1** Each 32-byte contains 4 words, thus every fourth access will occur one miss, the miss rate is: $\frac{1}{4}$

The miss rate is not sensitive to the size of the cache or the size of the working set.

All misses are compulsory miss.

- 5.8.2** When the cache block size is 16 bytes, each 16-bytes contains 2 word to store data. Thus every twice access will occur a miss.

The miss rate is: $\frac{1}{2}$

The same as above, the miss rate of the 64-bytes block cache is: $\frac{1}{8}$

The miss rate of 128-bytes is: $\frac{1}{16}$

5.8.3 In this case without including the compulsory miss on the first access, the miss rate is 0.

The pre-fetch buffer always has the next request ready

Exercise 5.9 Cache block size (B) can affect both miss rate and miss latency. Assuming a 1-CPI machine with an average of 1.35 references (both instruction and data) per instruction, help find the optimal block size given the following miss rates for various block sizes.

| | | | | |
|-------|--------|--------|----------|---------|
| 8: 4% | 16: 3% | 32: 2% | 64: 1.5% | 128: 1% |
|-------|--------|--------|----------|---------|

5.9.1 What is the optimal block size for a miss latency of $20 \times B$ cycles?

5.9.2 What is the optimal block size for a miss latency of $24 + B$ cycles?

5.9.3 For constant miss latency, what is the optimal block size?

SOLUTION:

5.9.1 AMAT for B=8: $0.04 \times (20 \times 8) = 6.4$
AMAT for B=16: $0.03 \times 20 \times 16 = 9.6$
AMAT for B=32: $0.02 \times 20 \times 32 = 12.8$
AMAT for B=64: $0.015 \times 20 \times 64 = 19.2$
AMAT for B=128: $0.01 \times 20 \times 128 = 25.6$
B=8 is optimal

5.9.2 AMAT for B=8: $0.04 \times (24 + 8) = 1.28$
AMAT for B=16: $0.03 \times (24 + 16) = 1.2$
AMAT for B=32: $0.02 \times (24 + 32) = 1.12$
AMAT for B=64: $0.015 \times (24 + 64) = 1.32$
AMAT for B=128: $0.01 \times (24 + 128) = 1.52$
Thus, B=32 is optimal

5.9.3 Because miss latency is constant, the optimal is B=128, the miss rate is the lowest.

Exercise 5.10 In this exercise, we will look at the different ways capacity affects overall performance. In general, cache access time is proportional to capacity. Assume that main memory accesses take 70ns and that memory accesses are 36% of all

instructions. The following table shows data for L1 caches attached to each of two processors, P1 and P2.

| | L1 Size | L1 Miss Rate | L1 Hit Time |
|----|---------|--------------|-------------|
| P1 | 2 KB | 8.0% | 0.66 ns |
| P2 | 4 KB | 6.0% | 0.90 ns |

- 5.10.1** Assuming that the L1 hit time determines the cycle times for P1 and P2, what are their respective clock rates?
- 5.10.2** What is the Average Memory Access Time for P1 and P2?
- 5.10.3** Assuming a base CPI of 1.0 without any memory stalls, what is the total CPI for P1 and P2? Which processor is faster? (When we say a “base CPI of 1.0”, we mean that instructions complete in one cycle, unless either the instruction access or the data access causes a cache miss.) For the next three problems, we will consider the addition of an L2 cache to P1 to presumably make up for its limited L1 cache capacity. Use the L1 cache capacities and hit times from the previous table when solving these problems. The L2 miss rate indicated is its local miss rate.

| | L2 Size | L2 Miss Rate | L2 Hit Time |
|--|---------|--------------|-------------|
| | 1 MB | 95% | 5.62 ns |

- 5.10.4** What is the AMAT for P1 with the addition of an L2 cache? Is the AMAT better Or worse with the L2 cache
- 5.10.5** Assuming a base CPI of 1.0 without any memory stalls, what is the total CPI for P1 with the addition of an L2 cache?
- 5.10.6** What would the L2 miss rate need to be in order for P1 with an L2 cache to be faster than P1 without an L2 cache?
- 5.10.7** What would the L2 miss rate need to be in order for P1 with an L2 cache to be faster than P2 without an L2 cache?

SOLUTION:

- 5.10.1** The hit time determines the clock cycle time of p1 and p2,

$$\text{And clock rate} = \frac{1}{\text{clock cycle time}}$$

$$\text{Thus, the clock rate of p1} = \frac{1}{0.66\text{ns}} = \frac{1}{0.66 \times 10^{-9}} = 1.515 \times 10^9 \text{Hz} = 1.515 \text{GHz}$$

$$\text{The clock rate of p2} = \frac{1}{0.9\text{ns}} = \frac{1}{0.9 \times 10^{-9}} = 1.111 \times 10^9 \text{Hz} = 1.111 \text{GHz}$$

- 5.10.2** AMAT of P1: AMAT of P1 = p1 hit time+ L1 miss rate×miss penalty
 AMAT of P1 = 0.66ns+0.08×70
 AMAT of P1 = 6.26ns

$$\begin{aligned} \text{AMAT of P2: AMAT of P2} &= \text{p2 hit time+ L1 miss rate} \times \text{miss penalty} \\ &= 0.9\text{ns} + 0.06 \times 70 \\ &= 5.1\text{ns} \end{aligned}$$

5.10.3 For P1, every instruction require one cycle, additionally, 8% of data access will be missed, thus, it require additional $0.08 \times (70/0.66) \approx 9$ cycles
 Furthermore, 36% of the instructions are data accesses. 8% of these 36% are cache misses, which adds an additional 107 cycles.
 The total CPI is:

$$1 + 9 + 0.08 \times 0.36 \times 107 = 12.96 \text{ cycles}$$

For P2, every instruction require on cycle, additionally, 8% of data access will be missed, thus, it require additional $0.06 \times (70/0.9) \approx 5$ cycles
 Furthermore, 36% of the instructions are data accesses. 8% of these 36% are cache misses, which adds an additional 107 cycles.
 The total CPI is:

$$1 + 5 + 0.06 \times 0.36 \times (70 \div 0.9) \approx 7.68 \text{ cycles}$$

5.10.4 For each L1 access, it will take 0.66ns, and for L1 miss, it will take access to L2 to find data, it will take 5.62ns, if L2 also miss, it will access main memory.
 Thus, the AMAT of P1 is: $0.66 + 0.08 \times (5.62 + 0.95 \times 70) \approx 6.42\text{ns}$
 It will be worse with L2.

5.10.5 One access for L1 requires one cycle, 8% of memory miss and access L2, it will require extra

$$5.62 \div 0.66 \approx 8.5\text{cycles}$$

And 95% of memory in L2 miss and access main memory, and 36% of instruction will access main memory.

Thus, the total CIP of P1 is:

$$1 + 0.08 \times (9 + 0.95 \times (70 \div 0.66)) + 0.36 \times 0.08 \times (9 + 0.95 \times (70 \div 0.66)) \approx 13\text{cycles}$$

5.10.6 Assume AMAT of P2 without of L2 < AMAT of P1 with L2

$$6.26 < 0.66 + 0.08 \times (5.62 + x \times 70)$$

$$x > 0.919$$

Thus, the highest miss rate of L2 is 92% if P1 with an L2 cache to be faster than P1 without an L2 cache.

5.10.7 Assume AMAT of P1 with of L2 < AMAT of P1 with L2

$$5.1 < 0.66 + 0.08 \times (5.62 + x \times 70)$$

$$x > 0.71$$

Thus, the highest miss rate of L2 is 71% if P1 with an L2 cache to be faster than P1 without an L2 cache.

Exercise 5.12 Multilevel caching is an important technique to overcome the limited amount of space that a first level cache can provide while still maintaining its speed.
 Consider a processor with the following parameters:

| Base CPI, No Memory Stalls | Processor Speed | Main Memory Access Time | First Level Cache MissRate per Instruction | Second Level Cache, Direct- Mapped Speed | Global Miss Rate with Second Level Cache, Direct- Mapped | Second Level Cache, Eight-Way Set Associative Speed | Global Miss Rate with Second Level Cache, Eight-Way Set Associative |
|-------------------------------------|--------------------|----------------------------------|--|---|--|---|---|
| 1.5 | 2 GHz | 100 ns | 7% | 12 cycles | 3.5% | 28 cycles | 1.5% |

**First Level Cache miss rate is per instruction. Assume the total number of L1 cache misses (instruction and data combined) is equal to 7% of the number of instructions.

- 5.12.1** Calculate the CPI for the processor in the table using: 1) only a first level cache, 2) a second level direct-mapped cache, and 3) a second level eight-way set associative cache. How do these numbers change if main memory access time is doubled? (Give each change as both an absolute CPI and a percent change.) Notice the extent to which an L2 cache can hide the effects of a slow memory.
- 5.12.2** It is possible to have an even greater cache hierarchy than two levels. Given the processor above with a second level, direct-mapped cache, a designer wants to add a third level cache that takes 50 cycles to access and will have a 13% miss rate. Would this provide better performance? In general, what are the advantages and disadvantages of adding a third level cache?

SOLUTION:

- 5.12.1** 1) For each access will require 1.5 cycle, 7% of access will miss, the extra cycle require is: $0.07 \times (100 \div 0.5) = 14$
Thus, the total time access cost is:
 $1.5 + 14 = 15.5$ cycles
- 2) With second cache, for each access will require 1.5 cycles, and 7% access will miss and access cache second level cache.
Thus, the total CPI is:
 $1.5 + 0.07 \times (12 + 0.035 \times (100 \div 0.5)) = 2.83$ cycles
- 3) For the 8-way set associated case, the total CPI is:
 $1.5 + 0.07 \times (28 + 0.0015 \times 200) = 3.67$

For the main memory access time is doubled

- 1) The total CPI of single cache is: $1.5 + 0.07 \times (200 \div 0.5) = 29.5$
- 2) The total CPI of a second level direct-mapped cache is:
 $1.5 + 0.07 \times (12 + 0.035 \times (200 \div 0.5)) = 3.32$ cycles
- 3) The total CPI of second level 8-ways set is:
 $1.5 + 0.07 \times (28 + 0.035 \times (200 \div 0.5)) = 3.88$ cycles

- 5.12.2** The total CPI of the third level cache is:

$$1.5 + 0.07 \times (12 + 0.035 \times (50 + 0.13 \times 200)) = 2.52$$
cycles

For the second level cache is 2.83cycles, the performance is better.

Exercise 5.19 The following table shows the contents of a 4-entry TLB.

| Entry-ID | Valid | VA Page | Modified | Protection | PA Page |
|----------|-------|---------|----------|------------|---------|
| 1 | 1 | 140 | 1 | RW | 30 |
| 2 | 0 | 40 | 0 | RX | 34 |
| 3 | 1 | 200 | 1 | RO | 32 |
| 4 | 1 | 280 | 0 | RW | 31 |

- 5.19.1** Under what scenarios would entry 2's valid bit be set to zero?
- 5.19.2** What happens when an instruction writes to VA page 30? When would a software managed TLB be faster than a hardware managed TLB?
- 5.19.3** What happens when an instruction writes to VA page 200?

SOLUTION:

- 5.19.1** It would be invalid if it was paged out to disk
- 5.19.2** A write to VA page 30 will occur a miss. When software can pre-fetch TLB entry, it will be faster than hardware.
- 5.19.3** An interrupt would be generated because the page is marked as read only.

Exercise 5.20 In this exercise, we will examine how replacement policies impact miss rate. Assume a 2-way set associative cache with 4 one word blocks. Consider the following word address sequence: 0, 1, 2, 3, 4, 2, 3, 4, 5, 6, 7, 0, 1, 2, 3, 4, 5, 6, 7, 0.

- 5.20.1** Assuming an LRU replacement policy, which accesses are hits?
- 5.20.2** Assuming an MRU (most recently used) replacement policy, which accesses are hits?
- 5.20.3** Simulate a random replacement policy by flipping a coin. For example, "heads" means to evict the first block in a set and "tails" means to evict the second block in a set. How many hits does this address sequence exhibit?
- 5.20.4** Describe an optimal replacement policy for this sequence. Which accesses are hits using this policy?
- 5.20.5** Describe why it is difficult to implement a cache replacement policy that is optimal for all address sequences.

5.20.6 Assume you could make a decision upon each memory reference whether or not you want the requested address to be cached. What impact could this have on miss rate?

SOLUTION:

5.20.1

Highlight in yellow is the LRU block.

| Address | Index | Tag | Block1 | | Block2 | | Miss or hit |
|---------|-------|-----|--------|-------|--------|-------|-------------|
| 0 | 000 | 0 | 0 | Empty | Empty | Empty | Miss |
| 1 | 001 | 0 | 0 | 1 | Empty | Empty | Miss |
| 2 | 010 | 1 | 0 | 1 | 2 | Empty | Miss |
| 3 | 011 | 1 | 0 | 1 | 2 | 3 | Miss |
| 4 | 100 | 0 | 1 | 4 | 2 | 3 | Miss |
| 2 | 010 | 1 | 0 | 4 | 2 | 3 | Hit |
| 3 | 011 | 1 | 1 | 4 | 2 | 3 | Hit |
| 4 | 100 | 0 | 1 | 4 | 2 | 3 | Hit |
| 5 | 101 | 0 | 1 | 4 | 5 | 2 | Miss |
| 6 | 110 | 1 | 1 | 4 | 5 | 6 | Miss |
| 7 | 111 | 1 | 1 | 4 | 5 | 6 | Miss |
| 0 | 000 | 0 | 0 | 0 | 5 | 6 | Miss |
| 1 | 001 | 0 | 0 | 0 | 1 | 6 | Miss |
| 2 | 010 | 1 | 0 | 0 | 1 | 2 | Miss |
| 3 | 011 | 1 | 0 | 0 | 1 | 2 | Miss |
| 4 | 100 | 0 | 1 | 4 | 1 | 2 | Miss |
| 5 | 101 | 0 | 1 | 4 | 5 | 2 | Miss |
| 6 | 110 | 1 | 1 | 4 | 5 | 6 | Miss |
| 7 | 111 | 1 | 1 | 4 | 5 | 6 | Miss |
| 0 | 000 | 0 | 0 | 0 | 5 | 6 | Miss |

5.20.2

Highlight in yellow is the MRU block.

| Address | Index | Tag | Block1 | | Block2 | | Miss or hit |
|---------|-------|-----|--------|-------|--------|-------|-------------|
| 0 | 000 | 0 | 0 | Empty | Empty | Empty | Miss |
| 1 | 001 | 0 | 0 | 1 | Empty | Empty | Miss |
| 2 | 010 | 1 | 0 | 1 | 2 | Empty | Miss |
| 3 | 011 | 1 | 0 | 1 | 2 | 3 | Miss |
| 4 | 100 | 0 | 1 | 1 | 2 | 4 | Miss |
| 2 | 010 | 1 | 0 | 1 | 2 | 4 | Hit |
| 3 | 011 | 1 | 1 | 0 | 1 | 3 | Miss |
| 4 | 100 | 0 | 1 | 0 | 1 | 3 | Hit |
| 5 | 101 | 0 | 1 | 0 | 1 | 3 | Miss |
| 6 | 110 | 1 | 1 | 0 | 1 | 3 | Miss |
| 7 | 111 | 1 | 1 | 0 | 1 | 3 | Miss |
| 0 | 000 | 0 | 0 | 0 | 1 | 3 | Hit |
| 1 | 001 | 0 | 0 | 0 | 1 | 3 | Hit |
| 2 | 010 | 1 | 0 | 0 | 2 | 3 | Miss |

| | | | | | | | | |
|---|-----|---|---|---|---|---|---|------|
| 3 | 011 | 1 | 0 | 0 | 2 | 3 | 7 | Hit |
| 4 | 100 | 0 | 1 | 0 | 2 | 4 | 7 | Miss |
| 5 | 101 | 0 | 1 | 0 | 2 | 5 | 7 | Miss |
| 6 | 110 | 1 | 1 | 0 | 2 | 6 | 7 | Miss |
| 7 | 111 | 1 | 1 | 0 | 2 | 6 | 7 | Hit |
| 0 | 000 | 0 | 0 | 0 | 2 | 6 | 7 | Hit |

5.20.3 The following is a Simulation results execute by myself.“Number” evict the first block in a set and “flower” means to evict the second block in a set.

| Address | | Block1 | | Block2 | | Miss or hit |
|---------|--------|--------|-------|--------|-------|-------------|
| 0 | Number | 0 | Empty | Empty | Empty | Miss |
| 1 | Flower | 0 | Empty | 1 | Empty | Miss |
| 2 | Flower | 0 | Empty | 1 | 2 | Miss |
| 3 | Flower | 0 | Empty | 3 | 2 | Miss |
| 4 | Number | 0 | 4 | 3 | 2 | Miss |
| 2 | Flower | 0 | 4 | 3 | 2 | Hit |
| 3 | Number | 0 | 4 | 3 | 2 | Hit |
| 4 | Flower | 0 | 4 | 3 | 2 | Hit |
| 5 | Flower | 0 | 4 | 3 | 5 | Miss |
| 6 | Flower | 0 | 4 | 6 | 5 | Miss |
| 7 | Number | 7 | 4 | 6 | 5 | Miss |
| 0 | Flower | 7 | 4 | 6 | 0 | Miss |
| 1 | Number | 7 | 1 | 6 | 0 | Miss |
| 2 | Number | 2 | 1 | 6 | 0 | Miss |
| 3 | Flower | 2 | 1 | 3 | 0 | Miss |
| 4 | Flower | 2 | 1 | 3 | 4 | Miss |
| 5 | Flower | 2 | 1 | 5 | 4 | Miss |
| 6 | Number | 2 | 6 | 5 | 4 | Miss |
| 7 | Flower | 2 | 6 | 5 | 7 | Miss |
| 0 | Number | 0 | 6 | 5 | 7 | Miss |

The miss rate is $\frac{17}{20} \approx 0.85$

5.20.4 MRU is an optimal policy

5.20.5 The best block to evict is the one that will cause the fewest misses in the future. Unfortunately, a cache controller cannot know the future! Our best alternative is to make a good prediction.

5.20.6 If you knew that an address had limited temporal locality and would conflict with another block in the cache, choosing not to cache it could improve the miss rate. On the other hand, you could worsen the miss rate by choosing poorly which addresses to cache.